

Efficient L^2 projection on simplices using Bernstein polynomials

Jesse Chan

Contents

1	Introduction	1
2	Bernstein-Bezier bases	1
3	Fast L^2 projections with Bernstein polynomials	1
3.1	Collapsed-coordinate quadratures	1
3.2	Polynomial multiplication	2
3.3	Inversion of modally diagonal matrices	2
4	Numerical experiments	3
4.1	Expansion kernels	3
4.1.1	Polynomial multiplication	3
4.1.2	Collapsed-coordinate quadrature	3
4.2	Mass matrix inversion kernel	3
4.3	Monolithic kernel	3
5	Application to Weight-adjusted Discontinuous Galerkin (WADG) methods	3

Abstract

Alternative formula for mass inversion. GPU-accelerated versions of Ainsworth and Kirby Duffy transforms. Additional polynomial multiplication-based approach.

1 Introduction

[1]

2 Bernstein-Bezier bases

3 Fast L^2 projections with Bernstein polynomials

Involves three steps - evaluation in an enriched representation (either at quadrature points or in a higher degree polynomial basis), scaling by the weight, and projection down to polynomials of degree N .

3.1 Collapsed-coordinate quadratures

Fast evaluation at Duffy points [2, 3, 4].

Component-wise scaling by weight evaluated at collapsed-coordinate quadrature points.

3.2 Polynomial multiplication

For non-constant coefficients, DG requires being able to deal with polynomial multiplication and projection onto lower-dimensional subspaces. Multiplying polynomials together may be done using a discrete convolution and polynomial multiplication (J. Sanchez-Reyes 2003). The projection operator may be derived by noting that degree elevation operators are diagonal when transformed to a modal basis.

Rescaling by binomial coefficients results in the unscaled Bernstein basis. Polynomial multiplication is then equivalent to discrete convolution of the scaled binomial coefficients.

Quadrature-free strategy for nonlinear volume terms: polynomial multiplication + projection.

1. Polynomial multiplication of two BB basis functions representable as coefficient scaling, N_p scalar multiplications and storage of N_p coeffs, and another coefficient scaling.
2. To reduce local memory costs, process coeffs for fg over one or more $(d - 1)$ dimensional layers.
3. Store ids and load a triangular number of loads.

3.3 Inversion of modally diagonal matrices

Any modally diagonal matrix \mathbf{M} can be represented in the form

$$\mathbf{M}_N^{-1} \mathbf{M}_{N,M} = \sum_{j=0}^N c_j E_{N-j}^N (E_{N-j}^M)^T.$$

This property was shown for the polynomial projection matrix by Waldron in [5, 6]. We give an alternative proof of this below where \mathbf{M} is any modally diagonal matrix.

The constants c_j may be computed through the solution of an $(N + 1) \times (N + 1)$ matrix system, using the fact that upon transformation to a modal basis, E_{N-j}^N is a diagonal matrix of ones and zeros, while E_{N-j}^M is a diagonal matrix with entries

$$\frac{\lambda_i^{N-j}}{\lambda_i^M}, \quad i = 0, \dots, N.$$

This may be factored into an application of E_N^M , then an application of

$$\begin{aligned} \sum_{j=0}^N c_j E_{N-j}^N (E_{N-j}^M)^T &= c_0 \mathbf{I} + c_1 E_{N-1}^N (E_{N-1}^M)^T + c_2 E_{N-1}^N E_{N-2}^{N-1} (E_{N-2}^{N-1})^T (E_{N-1}^M)^T + \dots \\ &= c_0 \mathbf{I} + c_1 E_{N-1}^N \left(\mathbf{I} + \frac{c_2}{c_1} E_{N-2}^{N-1} (\mathbf{I} + \dots) (E_{N-2}^{N-1})^T \right) (E_{N-1}^M)^T. \end{aligned}$$

This may be applied in two sweeps of length N , using in-place updates to memory. Unfortunately, for shared-memory parallelization, this will require synchronizations between each application of each matrix.

The cost of applying $(E_N^M)^T$ is the application of $(M - N)$ sparse degree elevation operations, each of which is $O(M^d)$ cost. Assuming $M \approx N$ (it is reasonable to match the order of the data with the order of approximation), this gives $O(N^{d+1})$ cost.

When applying the projection operator, since each operation is $O(N^3)$ and we apply $O(N)$ total operations, we have an $O(N^{d+1})$ overall cost.

This can also be used to apply the inverse mass matrix since it is diagonal under the transformation T .

4 Numerical experiments

4.1 Expansion kernels

4.1.1 Polynomial multiplication

4.1.2 Collapsed-coordinate quadrature

4.2 Mass matrix inversion kernel

4.3 Monolithic kernel

5 Application to Weight-adjusted Discontinuous Galerkin (WADG) methods

Numerical experiment: time to solution for order 7 and 8 using both NDG-WADG and BB-WADG. [7, 8]

References

- [1] Jesse Chan and T Warburton. GPU-accelerated Bernstein-Bezier discontinuous Galerkin methods for wave problems. *arXiv preprint arXiv:1512.06025*, 2015.
- [2] Mark Ainsworth, Miangaly Gaelle Andriamaro, and Oleg Davydov. Bernstein-Bézier finite elements of arbitrary order and optimal assembly procedures. *SIAM Journal on Scientific Computing*, 33(6):3087–3109, 2011.
- [3] Robert C Kirby. Fast simplicial finite element algorithms using Bernstein polynomials. *Numerische Mathematik*, 117(4):631–652, 2011.
- [4] Robert C Kirby and Kieu Tri Thinh. Fast simplicial quadrature-based finite element operators using Bernstein polynomials. *Numerische Mathematik*, 121(2):261–279, 2012.
- [5] Shayne Waldron. On the Bernstein-Bézier form of Jacobi polynomials on a simplex. *Journal of Approximation Theory*, 140(1):86–99, 2006.
- [6] Shayne Waldron. Computing orthogonal polynomials on a triangle by degree raising. *Numerical Algorithms*, 42(2):171–179, 2006.
- [7] Jesse Chan, Russell J Hewett, and T Warburton. Weight-adjusted discontinuous Galerkin methods: wave propagation in heterogeneous media. *arXiv preprint arXiv:1608.01944*, 2016.
- [8] Jesse Chan, Russell J Hewett, and T Warburton. Weight-adjusted discontinuous Galerkin methods: curvilinear meshes. *arXiv preprint arXiv:1608.03836*, 2016.