

Time-domain Bernstein-Bézier DG methods on simplices

Jesse Chan, T. Warburton

¹Department of Computational and Applied Mathematics
Rice University

27th Biennial Numerical analysis conference
University of Strathclyde
June 27, 2017

Outline and acknowledgments

Collaborators and contributors:

- T. Warburton (Virginia Tech)
- Russell J. Hewett (TOTAL E&P Research and Technology USA)

- 1 High order nodal DG methods
- 2 High order Bernstein-Bézier DG methods
- 3 Weight-adjusted DG: beyond low order coefficients/geometry

High order DG methods for wave propagation

- Unstructured (tetrahedral) meshes for geometric flexibility.
- Low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- High performance on many-core (explicit time stepping).

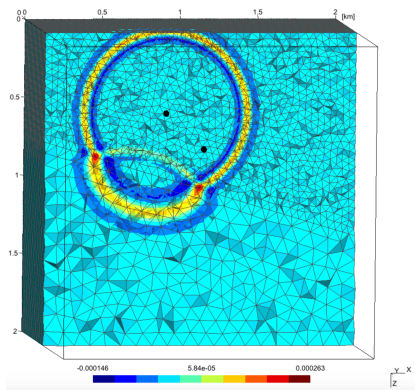
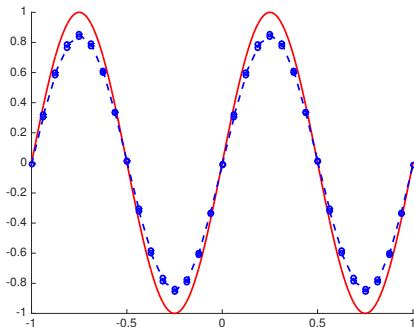


Figure courtesy of Axel Modave.

High order DG methods for wave propagation

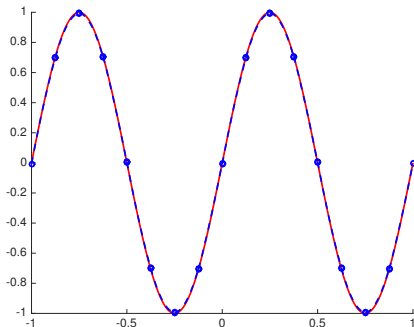
- Unstructured (tetrahedral) meshes for geometric flexibility.
- Low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- High performance on many-core (explicit time stepping).



Fine linear approximation.

High order DG methods for wave propagation

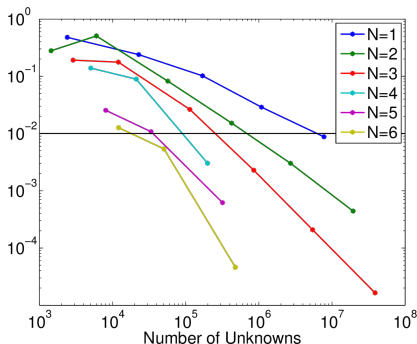
- Unstructured (tetrahedral) meshes for geometric flexibility.
- Low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- High performance on many-core (explicit time stepping).



Coarse quadratic approximation.

High order DG methods for wave propagation

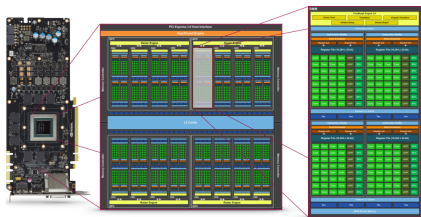
- Unstructured (tetrahedral) meshes for geometric flexibility.
- Low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- High performance on many-core (explicit time stepping).



Max errors vs. dofs.

High order DG methods for wave propagation

- Unstructured (tetrahedral) meshes for geometric flexibility.
- Low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- High performance on many-core (explicit time stepping).

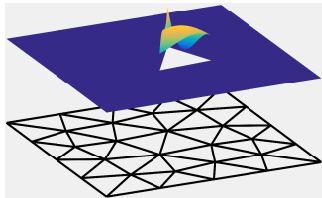


A graphics processing unit (GPU).

Discontinuous Galerkin methods

Discontinuous Galerkin (DG) methods:

- Piecewise polynomial approximation.
- Weak continuity across faces.
- Continuous PDE (for illustration: constant advection)



$$\frac{\partial u}{\partial t} = \frac{\partial u}{\partial x}$$

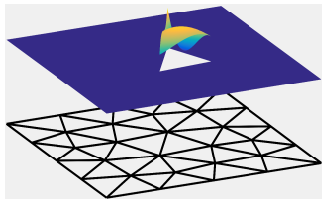
- DG local weak form over D_k with numerical flux \mathbf{f}^* .

$$\int_{D_k} \frac{\partial u}{\partial t} \phi = \int_{D_k} \frac{\partial u}{\partial x} \phi + \int_{\partial D_k} \mathbf{n} \cdot (\mathbf{f}^* - \mathbf{f}(u)) \phi, \quad u, \phi \in V_h$$

Discontinuous Galerkin methods

Discontinuous Galerkin (DG) methods:

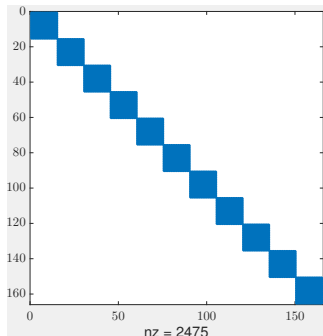
- Piecewise polynomial approximation.
- Weak continuity across faces.



DG yields system of ODEs

$$\mathbf{M}_{\Omega} \frac{d\mathbf{u}}{dt} = \mathbf{A}\mathbf{u}.$$

DG mass matrix decouples across elements, inter-element coupling only through \mathbf{A} .



Discontinuous Galerkin methods

Discontinuous Galerkin (DG) methods:

- Piecewise polynomial approximation.
- Weak continuity across faces.
- Matrix-free evaluation of $\mathbf{M}^{-1}\mathbf{A}$.
- Local differentiation and lifting matrices \mathbf{D}_x and $\mathbf{L}_f = \mathbf{M}^{-1}\mathbf{M}_f$.
- Assume (for now) piecewise constant coefficients and affine mappings.

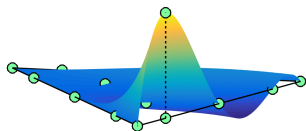
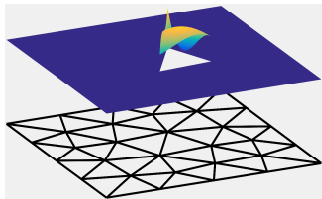
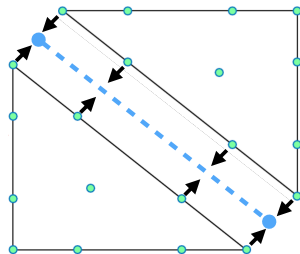


Figure: Nodal bases simplify flux computations.

Time-domain nodal DG methods

Given initial condition $u(\mathbf{x}, 0)$:

- Compute numerical flux at face nodes (**non-local**).
- Compute RHS of (**local**) ODE.
- Evolve (**local**) solution using explicit time integration (RK, AB, etc).



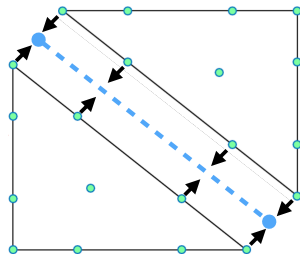
$$\frac{d\mathbf{u}}{dt} = \mathbf{D}_x \mathbf{u} + \sum_{\text{faces}} \mathbf{L}_f (\text{flux}), \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

Well-suited to GPUs computing (**significant** speedups over CPUs)!

Time-domain nodal DG methods

Given initial condition $u(\mathbf{x}, 0)$:

- Compute numerical flux at face nodes (**non-local**).
- Compute RHS of (**local**) ODE.
- Evolve (**local**) solution using explicit time integration (RK, AB, etc).



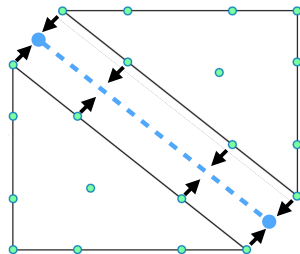
$$\frac{du}{dt} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f (\text{flux})}_{\text{Surface kernel}}, \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

Well-suited to GPUs computing (**significant** speedups over CPUs)!

Time-domain nodal DG methods

Given initial condition $u(\mathbf{x}, 0)$:

- Compute numerical flux at face nodes (**non-local**).
- Compute RHS of (**local**) ODE.
- Evolve (**local**) solution using explicit time integration (RK, AB, etc).



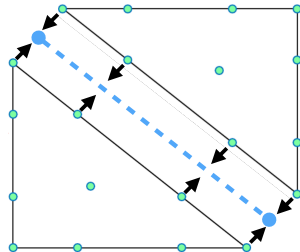
$$\underbrace{\frac{du}{dt}}_{\text{Update kernel}} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f (\text{flux})}_{\text{Surface kernel}}, \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

Well-suited to GPUs computing (**significant** speedups over CPUs)!

Time-domain nodal DG methods

Given initial condition $u(\mathbf{x}, 0)$:

- Compute numerical flux at face nodes (**non-local**).
- Compute RHS of (**local**) ODE.
- Evolve (**local**) solution using explicit time integration (RK, AB, etc).



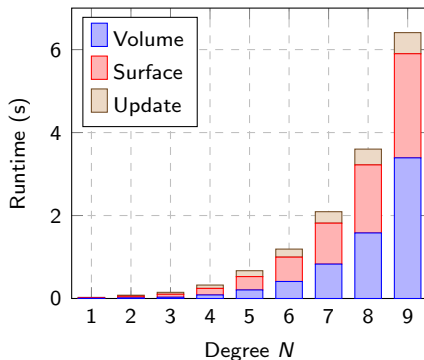
$$\underbrace{\frac{du}{dt}}_{\text{Update kernel}} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f (\text{flux})}_{\text{Surface kernel}}, \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

Well-suited to GPUs computing (**significant** speedups over CPUs)!

Computational costs at high orders of approximation

Problem: (tetrahedral) DG becomes **expensive** at high orders!

Nodal DG runtimes



- Large **dense** matrices:
 $O(N^6)$ work per tet.
- Tensor-product elements usually preferred for very high orders.
- $O(N^4)$ vs $O(N^6)$ cost, but less geometric flexibility.

DG runtimes for 50 timesteps, 98304 elements.

Spectral element methods

- Tensor product elements, Gauss-Legendre-Lobatto nodal basis.
- $O(N^{d+1})$ vs $O(N^{2d})$ work per element (differentiation, lifting).
- Hexahedral mesh generation more difficult.

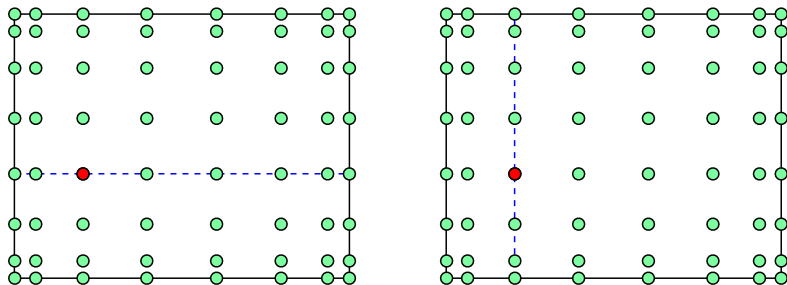


Figure: Spectral element stencils for $N = 7$ (orders $N > 10$ not uncommon!).

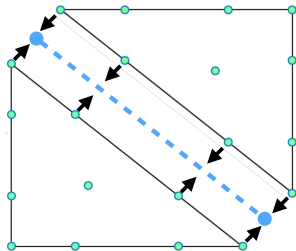
Fischer, Ronquist 1994. Spectral element methods for large scale parallel Navier-Stokes calculations.

Shepherd and Johnson 2008. Hexahedral mesh generation constraints.

High order nodal DG on tetrahedral meshes

$$\frac{d\mathbf{u}}{dt} = \mathbf{D}_x \mathbf{u} + \sum_{\text{faces}} \mathbf{L}_f (\text{flux}), \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

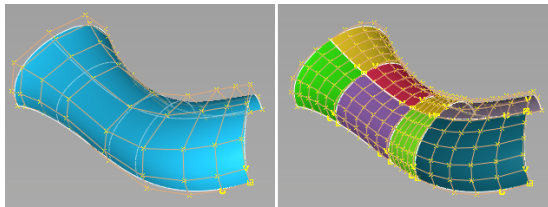
- Nodal bases: reduce the cost of computing numerical fluxes.
- No clear tetrahedral equivalent to spectral differentiation, lift matrices.
- $O(N^3)$ unknowns in 3D; $O(N^6)$ costs for applying **dense** matrices.



Derivative and lift matrices depend on the basis:
can we choose one that is efficient (and numerically stable)?

Bernstein-Bézier bases for finite element methods

- Geometry, graphics, Computer Aided Design (CAD).



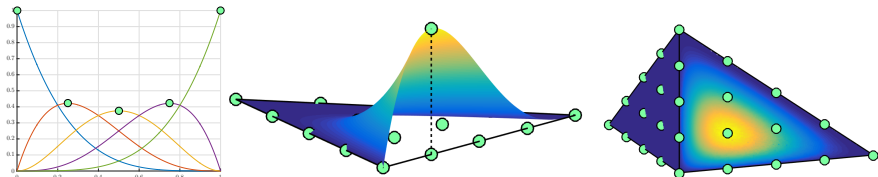
- Recent developments: optimal complexity assembly of finite element matrices, sum factorization (reduced complexity quadrature).
- This work: adapt Bernstein-Bézier for time-domain DG methods.

Split multi-span NURBS surfaces into Bézier patches, <https://knowledge.autodesk.com>

Ainsworth et al. 2011. Bernstein-Bézier finite elements of arbitrary order and optimal assembly procedures.

Kirby 2011. Fast simplicial finite element algorithms using Bernstein polynomials.

Bernstein-Bézier polynomial bases on simplices



Each function attains its maximum at an equispaced lattice point of a d -simplex.

■ Simple expression in 1D

$$B_i^N(x) = x^i(1-x)^{N-i}, \quad 0 \leq x \leq 1.$$

■ Barycentric monomials on a d -simplex. For a tetrahedron,

$$B_{ijkl}^N(\lambda_0, \lambda_1, \lambda_2, \lambda_3) = \frac{N!}{i!j!k!l!} \lambda_0^i \lambda_1^j \lambda_2^k \lambda_3^l, \quad i+j+k+l = N.$$

■ Similar structure to nodal basis (vertex, edge, face, interior functions).

Bernstein-Bézier derivatives and degree elevation in 1D

- Simple differentiation of Bernstein polynomials

$$\frac{\partial B_i^N(x)}{\partial x} = N \left(B_{i-1}^{N-1}(x) - B_i^{N-1}(x) \right).$$

- Simple degree elevation of Bernstein polynomials

$$B_i^{N-1}(x) = \left(\frac{N-i}{N} \right) B_i^N(x) - \left(\frac{i+1}{N} \right) B_{i+1}^N(x).$$

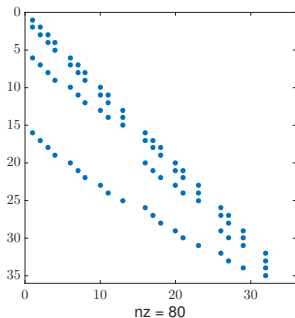
- Combine to get expansion of Bernstein derivatives

$$\frac{\partial B_i^N(x)}{\partial x} = a_i^N B_{i-1}^N(x) + b_i^N B_i^N(x) - c_i^N B_{i+1}^N(x).$$

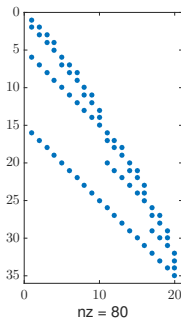
Implies 1D derivative matrix \mathbf{D}_x is **sparse** (tridiagonal).

Bernstein-Bézier derivative and degree elevation in 3D

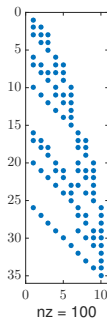
- Bernstein-Bézier barycentric differentiation matrices very sparse.
- Degree elevation matrices \mathbf{E}_{N-i}^N are sparse (for consecutive degrees).
- Higher degree elevation \rightarrow product of matrices $\mathbf{E}_{N-2}^N = \mathbf{E}_{N-1}^N \mathbf{E}_{N-2}^{N-1}$.



(a) Derivative matrix w.r.t. first barycentric coordinate.



(b) Deg. elevation matrix \mathbf{E}_{N-1}^N



(c) \mathbf{E}_{N-2}^N

Stencils for Bernstein-Bézier derivative matrices

- Stencil sizes at most $(d + 1)$ in d dimensions.
- Compute derivatives w.r.t. barycentric coordinates.
- Stencil values are identical for **all** barycentric derivatives.

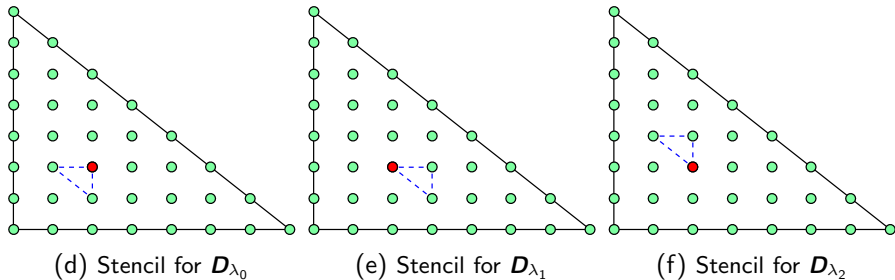
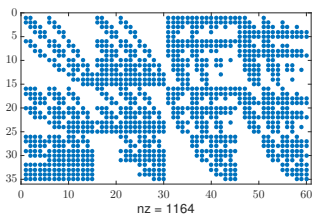


Figure: Bernstein-Bézier stencils for a single node (in red) $N = 7$.

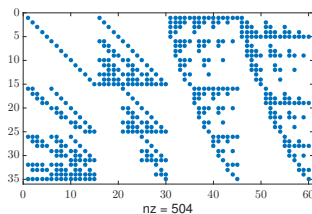
Factorization of the Bernstein lift operator

The Bernstein-Bézier lift matrix \mathbf{L} admits a factorization of the form

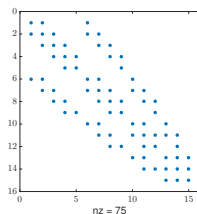
$$\mathbf{L} = \mathbf{E}_L \begin{pmatrix} \mathbf{L}_0 & & & \\ & \mathbf{L}_0 & & \\ & & \mathbf{L}_0 & \\ & & & \mathbf{L}_0 \end{pmatrix}.$$



(a) Lift matrix \mathbf{L}



(b) Lift reduction \mathbf{E}_L

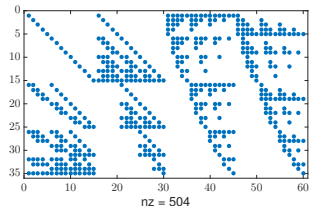


(c) \mathbf{L}_0

Factorization of the Bernstein lift operator

The Bernstein-Bézier lift matrix \mathbf{L} admits a factorization of the form

$$\mathbf{L} = \mathbf{E}_L \begin{pmatrix} \mathbf{L}_0 & & & \\ & \mathbf{L}_0 & & \\ & & \mathbf{L}_0 & \\ & & & \mathbf{L}_0 \end{pmatrix}.$$



$$\mathbf{E}_L^1 = \begin{bmatrix} \mathbf{I} \\ \ell_1 (\mathbf{E}_{N-1}^N)^T \\ \vdots \\ \ell_N (\mathbf{E}_0^N)^T \end{bmatrix}.$$

$$\mathbf{E}_L = [\mathbf{E}_L^1 \mid \dots \mid \mathbf{E}_L^4] \text{ (4 faces).}$$

2D **degree reduction** matrices $(\mathbf{E}_i^N)^T$.

Bernstein-Bézier lift matrix: optimal complexity application

- L “lifts” numerical fluxes from faces to volume.
- Apply L_0 to face flux, extend to each “layer” of the simplex.

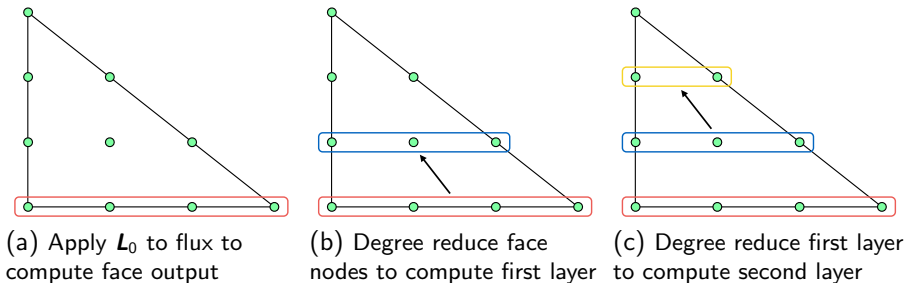


Figure: An $O(N^d)$ storage/complexity approach to applying the lift matrix.

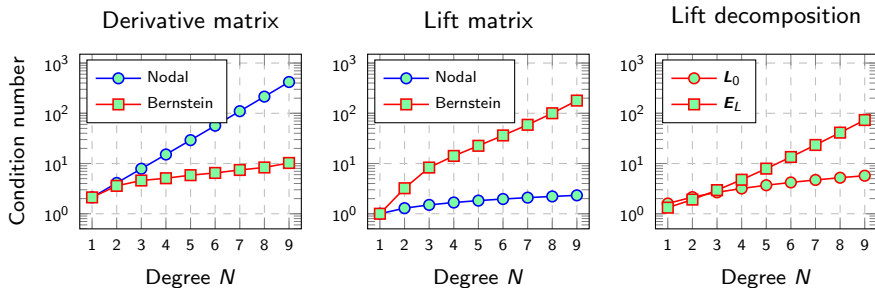
For $N < 6$, currently more efficient to treat E_L as a sparse matrix — irregular data accesses with optimal $O(N^d)$ approach.

Numerical stability of Bernstein-Bézier DG

- “Condition number” of Bernstein differentiation and lift matrices comparable to that of nodal bases.

$$\kappa(\mathbf{A}) = \frac{\sigma_1}{\sigma_r}$$

- Comparable long-time growth of (single precision) numerical error.



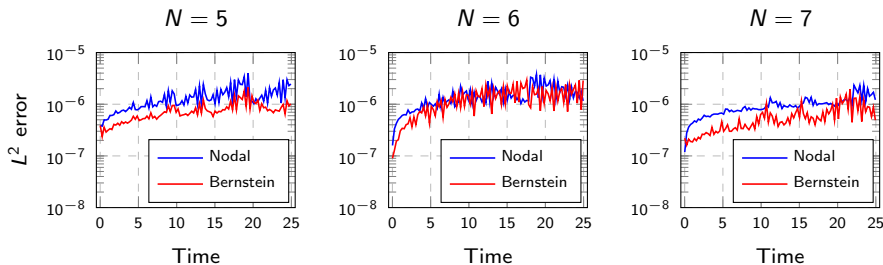
Condition numbers of matrices for nodal and Bernstein-Bézier bases.

Numerical stability of Bernstein-Bézier DG

- “Condition number” of Bernstein differentiation and lift matrices comparable to that of nodal bases.

$$\kappa(\mathbf{A}) = \frac{\sigma_1}{\sigma_r}$$

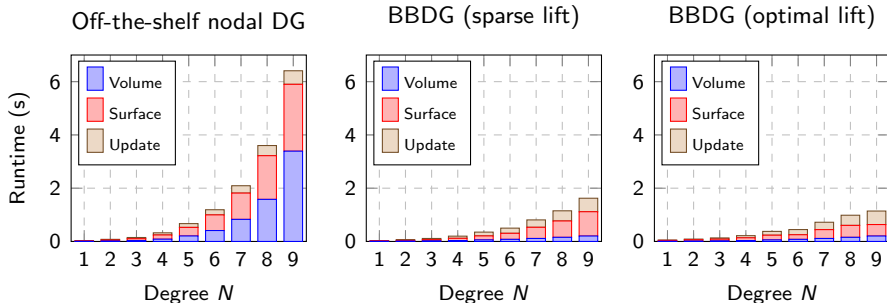
- Comparable long-time growth of (single precision) numerical error.



Evolution of L^2 error (acoustics) for nodal and Bernstein-Bézier bases.

GPU runtime comparisons of BBDG and nodal DG

Bernstein-Bézier DG achieves $\approx 2\times$ speedup at moderate orders, and up to $\approx 6\times$ speedup at high orders (for acoustics).

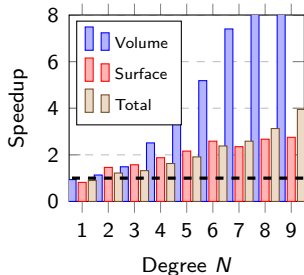


$$\underbrace{\frac{du}{dt}}_{\text{Update kernel}} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f (\text{flux})}_{\text{Surface kernel}}, \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

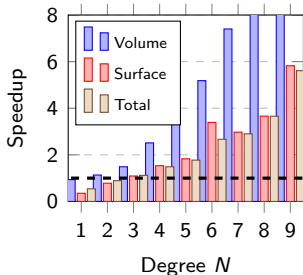
GPU runtime comparisons of BBDG and nodal DG

Bernstein-Bézier DG achieves $\approx 2\times$ speedup at moderate orders, and up to $\approx 6\times$ speedup at high orders (for acoustics).

BB speedup (sparse lift) over nodal



BB speedup (optimal lift) over nodal



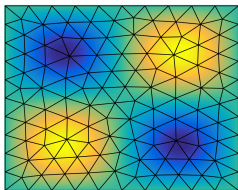
$$\underbrace{\frac{du}{dt}}_{\text{Update kernel}} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f (\text{flux})}_{\text{Surface kernel}}, \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

Extensions: high order models of heterogeneous media

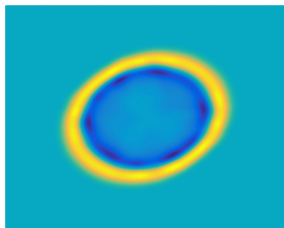
- Acoustic wave equation in heterogeneous media

$$\frac{1}{c^2(\mathbf{x})} \frac{\partial^2 p}{\partial t^2} - \Delta p = 0.$$

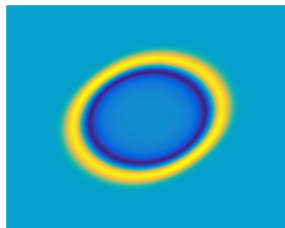
- Piecewise constant $c^2(\mathbf{x})$ efficient, but generates spurious reflections.
- Goal: high order $c^2(\mathbf{x})$, stability, low computational complexity.



(l) Mesh and exact c^2



(m) Piecewise constant c^2



(n) High order c^2

Weighted mass matrices and weight-adjusted DG

- Weighted mass matrix: high order accurate and energy stable, but high storage costs, $O(N^6)$ complexity to apply \mathbf{M}_w^{-1} .

$$\frac{d}{dt} \mathbf{M}_w \mathbf{u} = \text{right hand side}, \quad w = 1/c^2.$$

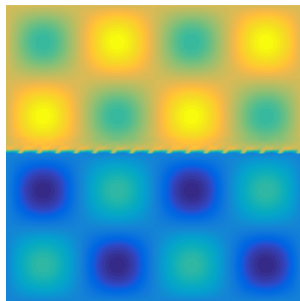
- **Weight-adjusted DG (WADG)**: energy stable, low storage approximation of weighted mass matrix

$$\frac{d}{dt} \mathbf{M}_w \mathbf{u} \approx \frac{d}{dt} \mathbf{M} (\mathbf{M}_{1/w})^{-1} \mathbf{M} \mathbf{u} = \text{right hand side}.$$

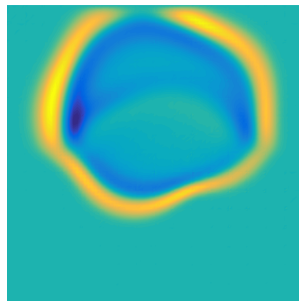
- Bypass inverse of weighted matrix $(\mathbf{M}_w)^{-1}$

$$\begin{aligned} \mathbf{M} (\mathbf{M}_{1/w})^{-1} \mathbf{M} \frac{d\mathbf{U}}{dt} &= \mathbf{A}_h \mathbf{U} \\ \rightarrow \frac{d\mathbf{U}}{dt} &= \mathbf{M}^{-1} \mathbf{M}_{1/w} \underbrace{\mathbf{M}^{-1} \mathbf{A}_h \mathbf{U}}_{\text{RHS for } w=1} \end{aligned}$$

Acoustic wave equation: heterogeneous media



(a) $c^2(x, y)$



(b) Standard DG

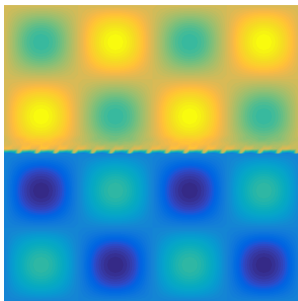
- L^2 convergence between optimal $O(h^{N+1})$, provable $O(h^{N+1/2})$.
- Extensions to curved elements, matrix weights (elastodynamics).

Chan, Hewett, Warburton. 2016. Weight-adjusted DG methods: wave propagation in heterogeneous media.

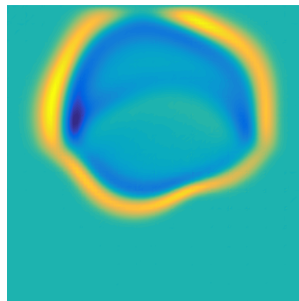
Chan, Hewett, Warburton. 2016. Weight-adjusted DG methods: curvilinear meshes.

Chan 2017. Weight-adjusted DG methods: matrix-valued weights and elastic wave prop. in heterogeneous media.

Acoustic wave equation: heterogeneous media



(a) $c^2(x, y)$



(b) Weighted-adjusted DG

- L^2 convergence between optimal $O(h^{N+1})$, provable $O(h^{N+1/2})$.
- Extensions to curved elements, matrix weights (elastodynamics).

Chan, Hewett, Warburton. 2016. Weight-adjusted DG methods: wave propagation in heterogeneous media.

Chan, Hewett, Warburton. 2016. Weight-adjusted DG methods: curvilinear meshes.

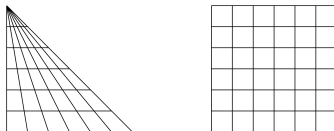
Chan 2017. Weight-adjusted DG methods: matrix-valued weights and elastic wave prop. in heterogeneous media.

WADG: low-complexity implementations

- Low storage, matrix-free application of $(\mathbf{M}_w^{-1} \mathbf{M})^{-1} = \mathbf{M}^{-1} \mathbf{M}_w$.

$$(\mathbf{M})^{-1} \mathbf{M}_{1/w} \text{RHS} = \underbrace{\widehat{\mathbf{M}}^{-1} \mathbf{V}_q^T W \text{diag}(1/w)}_{\mathbf{P}_q} \mathbf{V}_q (\text{RHS}).$$

- $O(N^4)$ cost in 3D: sum factorization for \mathbf{V}_q , block LDL for $\widehat{\mathbf{M}}^{-1}$.



- Current work: for **fixed** approximations of $w(\mathbf{x})$, **optimal** complexity WADG using polynomial multiplication and truncation.

Kirby 2017. Fast inversion of the simplicial Bernstein mass matrix.

Kirby and Tinh 2012. Fast simplicial quadrature-based finite element operators using Bernstein polynomials.

Summary and acknowledgements

- Optimal complexity RHS evaluation for time-domain DG.
- Bernstein-Bézier sparsity: efficiency at high orders on GPUs.

Thanks to NSF and TOTAL E&P Research and Technology USA
for their support of this work.

Chan 2017. Weight-adjusted DG methods: matrix-valued weights and elastic wave prop. in heterogeneous media (arXiv).

Chan, Hewett, Warburton. 2016. Weight-adjusted DG methods: wave propagation in heterogeneous media (SISC).

Chan, Hewett, Warburton. 2016. Weight-adjusted DG methods: curvilinear meshes (arXiv).

Chan, Warburton 2016. GPU-accelerated Bernstein-Bézier DG methods for wave problems (SISC).

Additional slides

Performance comparisons of BBDG and nodal DG

BBDG: lower FLOPs per second than nodal DG...
but maintains throughput/bandwidth as N increases!

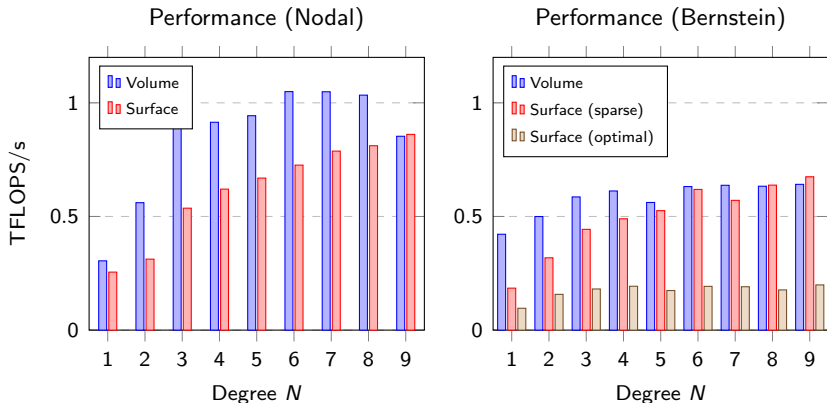


Figure: Profiled FLOPS/s for nodal and Bernstein-Bézier DG methods.

Performance comparisons of BBDG and nodal DG

BBDG: lower FLOPs per second than nodal DG...
but maintains throughput/bandwidth as N increases!

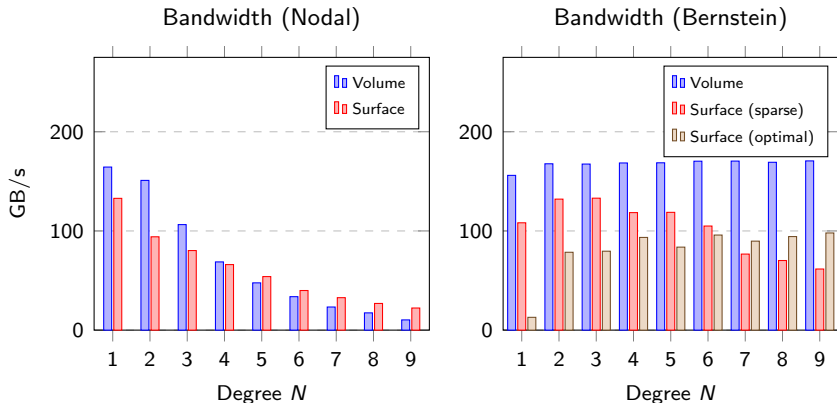
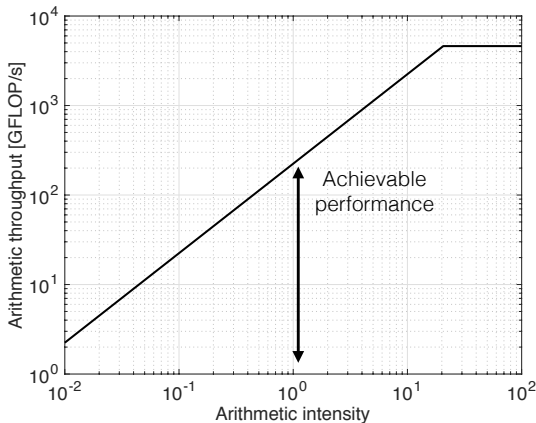


Figure: Profiled bandwidth for nodal and Bernstein-Bézier DG methods.

Roofline model: estimating computational efficiency

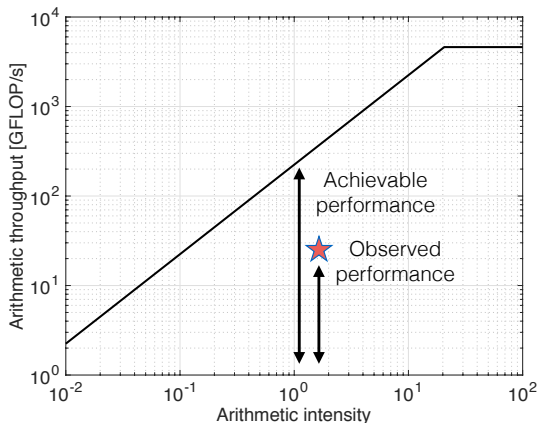
- Arithmetic intensity: floating-point operations per byte of data.
- Computational efficiency: ratio of observed/achievable performance.



Williams, Waterman, Patterson 2009. Roofline: an insightful visual performance model for multicore architectures.

Roofline model: estimating computational efficiency

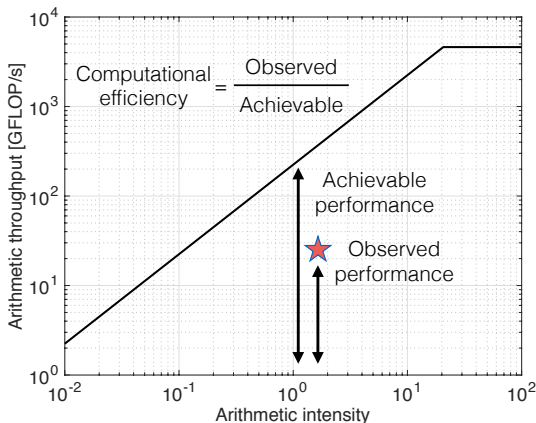
- Arithmetic intensity: floating-point operations per byte of data.
- Computational efficiency: ratio of observed/achievable performance.



Williams, Waterman, Patterson 2009. Roofline: an insightful visual performance model for multicore architectures.

Roofline model: estimating computational efficiency

- Arithmetic intensity: floating-point operations per byte of data.
- Computational efficiency: ratio of observed/achievable performance.



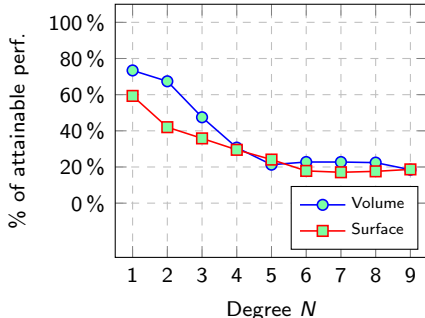
Williams, Waterman, Patterson 2009. Roofline: an insightful visual performance model for multicore architectures.

Efficiency comparisons of BBDG and nodal DG

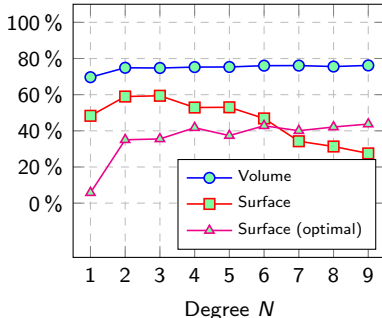
Bernstein-Bézier DG: standard implementation, sparse matrices.

$$\underbrace{\frac{d\mathbf{u}}{dt}}_{\text{Update kernel}} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f (\text{flux})}_{\text{Surface kernel}}, \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

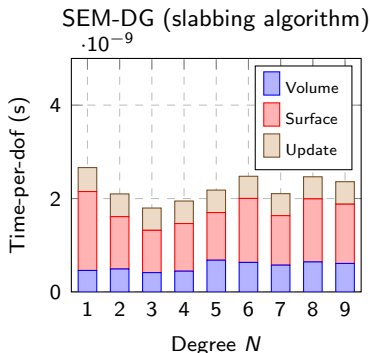
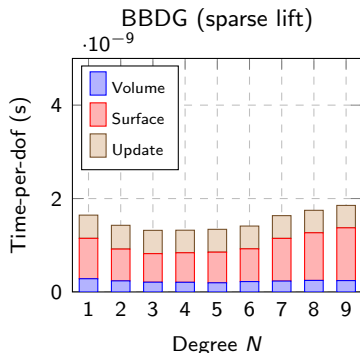
Efficiency (Standard)



Efficiency (Bernstein)

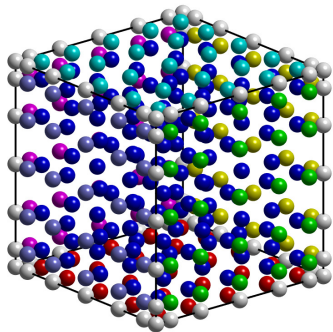


Runtime-only comparisons: BBDG, SEM-DG on GPUs



- BBDG 1-1.75 \times faster per dof than SEM-DG for $N \leq 10$.
- Unstructured hex meshes: $9(N+1)^3$ geometric factors per element.
- **Disclaimer:** hexes are more accurate, need time-to-error studies!

Runtime-only comparisons: BBDG, SEM-DG on GPUs



- BBDG $1\text{--}1.75\times$ faster per dof than SEM-DG for $N \leq 10$.
- Unstructured hex meshes: $9(N + 1)^3$ geometric factors per element.
- **Disclaimer:** hexes are more accurate, need time-to-error studies!