

Weight-adjusted Bernstein-Bezier DG methods for wave propagation in heterogeneous media

Jesse Chan, Kaihang Guo

Department of Computational and Applied Math, Rice University

Rice Oil and Gas HPC Conference 2018
March 12-13, 2018

High order DG methods for wave propagation

- Unstructured (tetrahedral) meshes for geometric flexibility.
- High order: low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- Explicit time stepping: high performance on many-core.

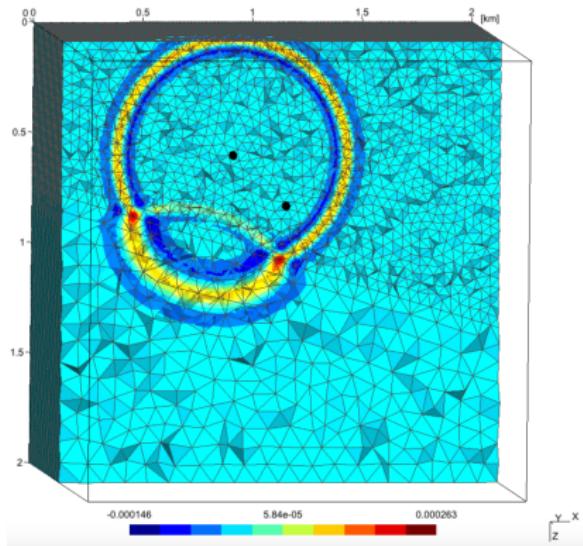
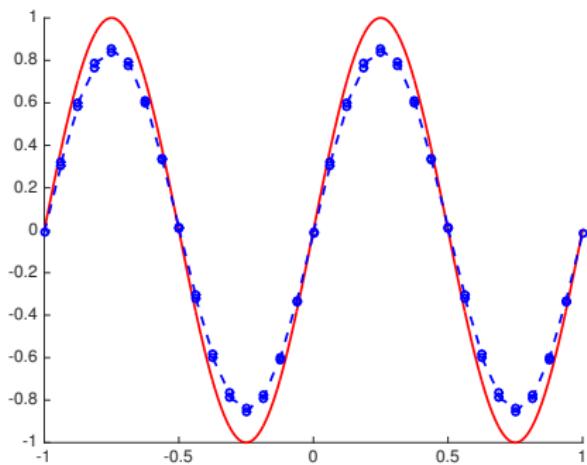


Figure courtesy of Axel Modave.

Goal: accuracy **and** efficiency for heterogeneous media.

High order DG methods for wave propagation

- Unstructured (tetrahedral) meshes for geometric flexibility.
- High order: low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- Explicit time stepping: high performance on many-core.

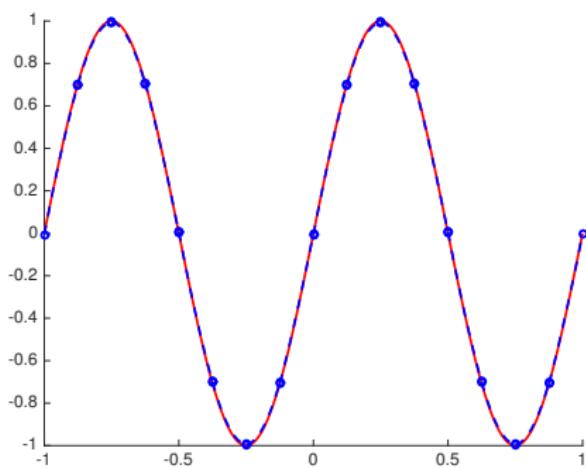


Fine linear approximation.

Goal: accuracy **and** efficiency for heterogeneous media.

High order DG methods for wave propagation

- Unstructured (tetrahedral) meshes for geometric flexibility.
- High order: low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- Explicit time stepping: high performance on many-core.

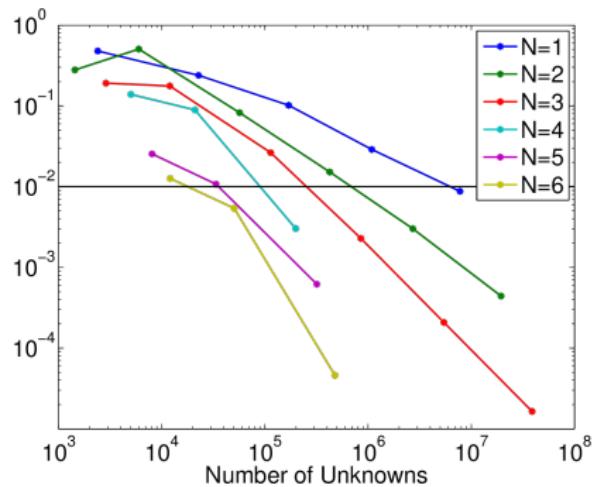


Coarse quadratic approximation.

Goal: accuracy **and** efficiency for heterogeneous media.

High order DG methods for wave propagation

- Unstructured (tetrahedral) meshes for geometric flexibility.
- High order: low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- Explicit time stepping: high performance on many-core.

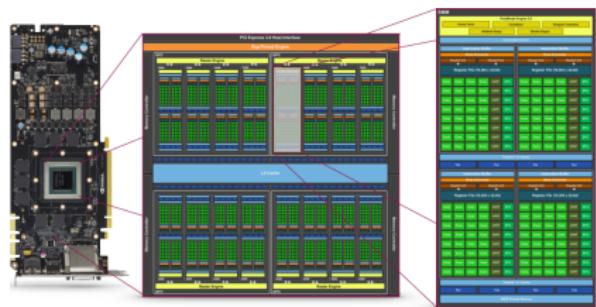


Max errors vs. dofs.

Goal: accuracy **and** efficiency for heterogeneous media.

High order DG methods for wave propagation

- Unstructured (tetrahedral) meshes for geometric flexibility.
- High order: low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- Explicit time stepping: high performance on many-core.

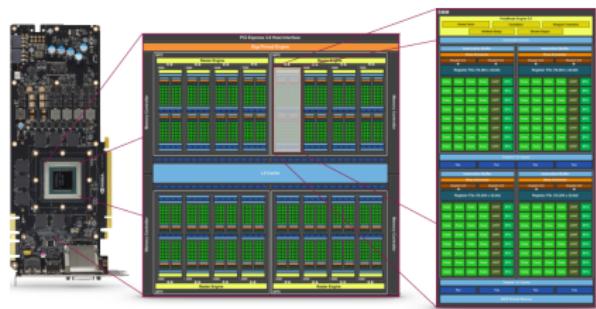


Graphics processing units (GPU).

Goal: accuracy **and** efficiency for heterogeneous media.

High order DG methods for wave propagation

- Unstructured (tetrahedral) meshes for geometric flexibility.
- High order: low numerical dissipation and dispersion.
- High order approximations: more accurate per unknown.
- Explicit time stepping: high performance on many-core.



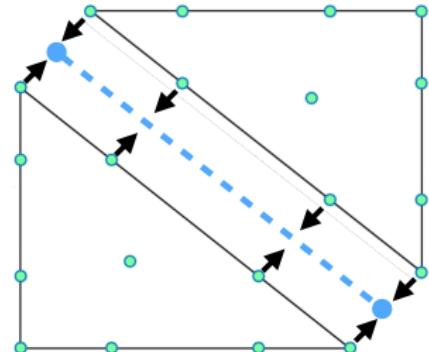
Graphics processing units (GPU).

Goal: accuracy **and** efficiency for heterogeneous media.

Time-domain nodal DG methods

Assume $u(\mathbf{x}, t) = \sum \mathbf{u}_j \phi_j(\mathbf{x})$ on D^k

- Compute numerical flux at face nodes (**non-local**).
- Compute RHS of (**local**) ODE.
- Evolve (**local**) solution using explicit time integration (RK, AB, etc).



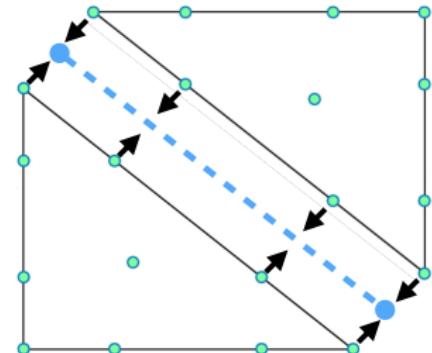
$$\frac{d\mathbf{u}}{dt} = \mathbf{D}_x \mathbf{u} + \sum_{\text{faces}} \mathbf{L}_f \text{ (flux)}.$$

$$\begin{aligned}\mathbf{M}_{ij} &= \int_{D^k} \phi_j(\mathbf{x}) \phi_i(\mathbf{x}) \\ \mathbf{L}_f &= \mathbf{M}^{-1} \mathbf{M}_f.\end{aligned}$$

Time-domain nodal DG methods

Assume $u(\mathbf{x}, t) = \sum \mathbf{u}_j \phi_j(\mathbf{x})$ on D^k

- Compute numerical flux at face nodes (**non-local**).
- Compute RHS of (**local**) ODE.
- Evolve (**local**) solution using explicit time integration (RK, AB, etc).



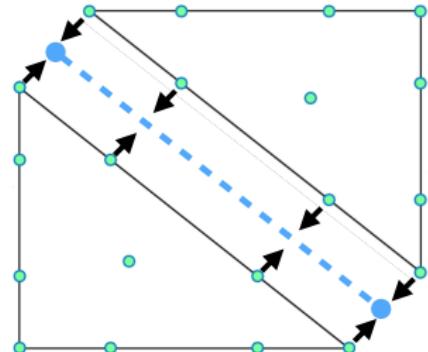
$$\frac{d\mathbf{u}}{dt} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f}_{\text{Surface kernel}} (\text{flux}).$$

$$\mathbf{M}_{ij} = \int_{D^k} \phi_j(\mathbf{x}) \phi_i(\mathbf{x})$$
$$\mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

Time-domain nodal DG methods

Assume $u(\mathbf{x}, t) = \sum \mathbf{u}_j \phi_j(\mathbf{x})$ on D^k

- Compute numerical flux at face nodes (**non-local**).
- Compute RHS of (**local**) ODE.
- Evolve (**local**) solution using explicit time integration (RK, AB, etc).



$$\underbrace{\frac{d\mathbf{u}}{dt}}_{\text{Update kernel}} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f (\text{flux})}_{\text{Surface kernel}}.$$

$$\mathbf{M}_{ij} = \int_{D^k} \phi_j(\mathbf{x}) \phi_i(\mathbf{x})$$
$$\mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

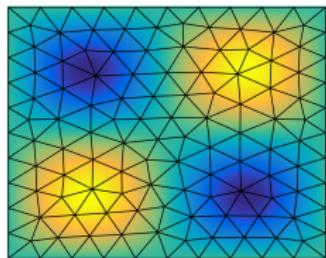
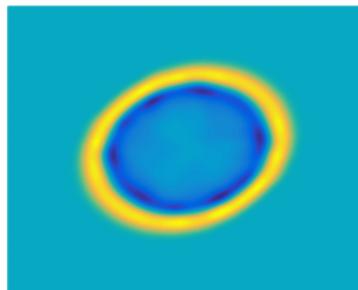
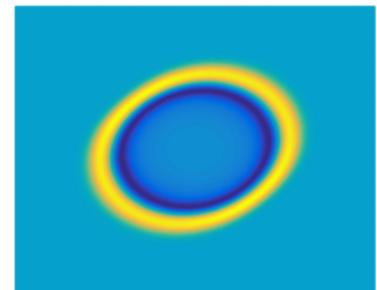
Outline

- 1 Weight-adjusted DG (WADG): arbitrary heterogeneous media
- 2 Bernstein-Bezier WADG: high order efficiency

Outline

- 1 Weight-adjusted DG (WADG): arbitrary heterogeneous media
- 2 Bernstein-Bezier WADG: high order efficiency

High order approximation of media and geometry

(a) Mesh and exact c^2 (b) Piecewise const. c^2 (c) High order c^2

- Piecewise constant wavespeed c^2 : efficient, but spurious reflections.

$$\frac{1}{c^2(\mathbf{x})} \frac{\partial p}{\partial t} + \nabla \cdot \mathbf{u} = 0, \quad \frac{\partial \mathbf{u}}{\partial t} + \nabla p = 0.$$

- High order wavespeeds: weighted mass matrices. Stable, but requires pre-computation/storage of inverses or factorizations!

$$\mathbf{M}_{1/c^2} \frac{d\mathbf{p}}{dt} = \mathbf{A}_h \mathbf{U}, \quad (\mathbf{M}_{1/c^2})_{ij} = \int_{D^k} \frac{1}{c^2(\mathbf{x})} \phi_j(\mathbf{x}) \phi_i(\mathbf{x}).$$

Weight-adjusted DG: stable, accurate, non-invasive

- Weight-adjusted DG (WADG): energy stable approx. of \mathbf{M}_{1/c^2}

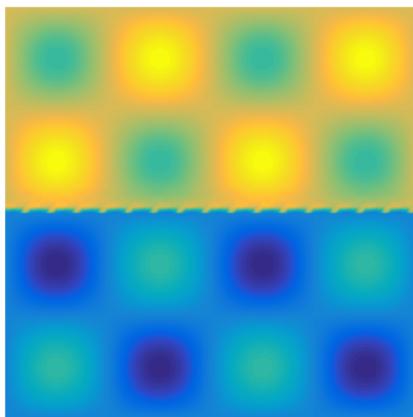
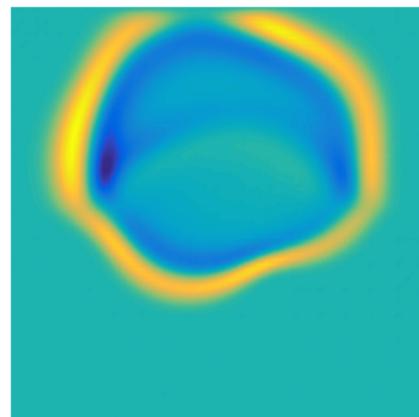
$$\mathbf{M}_{1/c^2} \frac{d\mathbf{p}}{dt} \approx \mathbf{M} (\mathbf{M}_{c^2})^{-1} \mathbf{M} \frac{d\mathbf{p}}{dt} = \mathbf{A}_h \mathbf{U}.$$

- New evaluation reuses implementation for constant wavespeed

$$\frac{d\mathbf{p}}{dt} = \underbrace{\mathbf{M}^{-1} (\mathbf{M}_{c^2})}_{\text{modified update}} \quad \underbrace{\mathbf{M}^{-1} \mathbf{A}_h \mathbf{U}}_{\text{constant wavespeed RHS}}$$

- Low storage matrix-free application of $\mathbf{M}^{-1} \mathbf{M}_{c^2}$ using quadrature-based interpolation and L^2 projection matrices $\mathbf{V}_q, \mathbf{P}_q$.

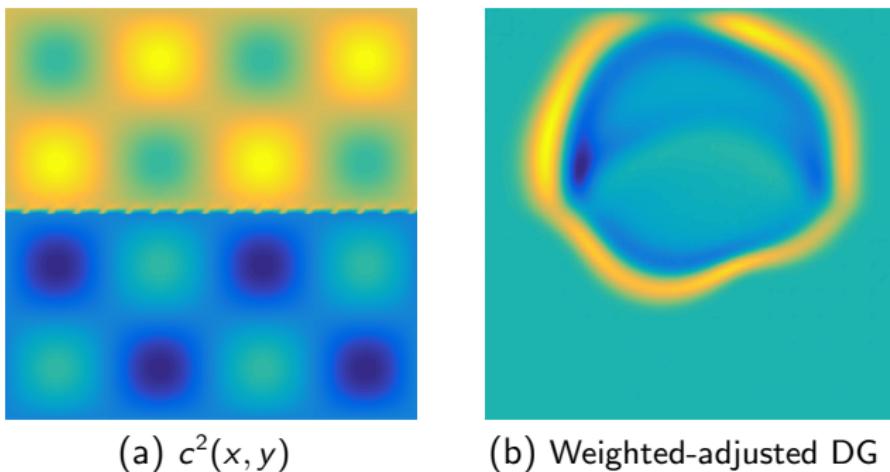
$$(\mathbf{M})^{-1} \mathbf{M}_{c^2} \text{RHS} = \underbrace{\mathbf{M}^{-1} \mathbf{V}_q^T W \text{diag}(c^2)}_{\mathbf{P}_q} \mathbf{V}_q \text{ (RHS).}$$

WADG: nearly identical to using M_{1/c^2}^{-1} (a) $c^2(x, y)$ 

(b) Standard DG

Figure: Standard vs. weight-adjusted DG with spatially varying c^2 .

- L^2 error is $O(h^{N+1})$; standard DG and WADG difference is $O(h^{N+2})$.
- Can generalize to matrix weights (elastic wave propagation).

WADG: nearly identical to using M_{1/c^2}^{-1} Figure: Standard vs. weight-adjusted DG with spatially varying c^2 .

- L^2 error is $O(h^{N+1})$; standard DG and WADG difference is $O(h^{N+2})$.
- Can generalize to matrix weights (elastic wave propagation).

WADG: more efficient than storing M_{1/c^2}^{-1} on GPUs

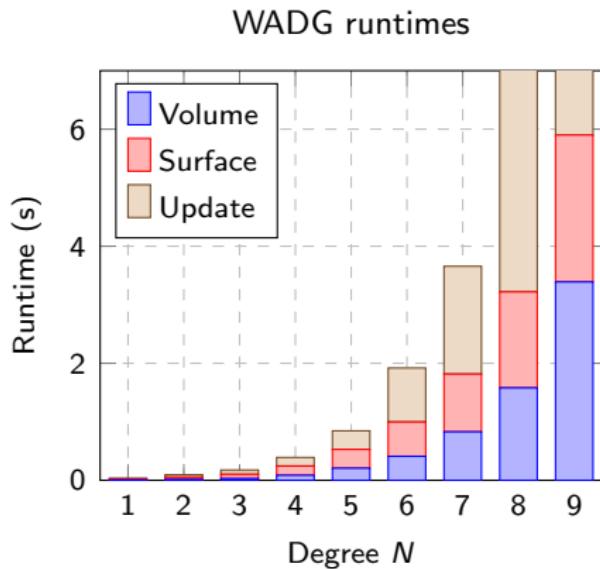
	$N = 1$	$N = 2$	$N = 3$	$N = 4$	$N = 5$	$N = 6$	$N = 7$
M_{1/c^2}^{-1}	.66	2.79	9.90	29.4	73.9	170.5	329.4
WADG	0.59	1.44	4.30	13.9	43.0	107.8	227.7
Speedup	1.11	1.94	2.30	2.16	1.72	1.58	1.45

Time (ns) per element: storing/applying M_{1/c^2}^{-1} vs WADG (deg. $2N$ quadrature).

- Efficiency on GPUs: reduce memory accesses and data movement.
- (Tuned) low storage WADG faster than storing and applying M_{1/c^2}^{-1} !

Computational costs at high orders of approximation

Problem: WADG at high orders becomes **expensive!**

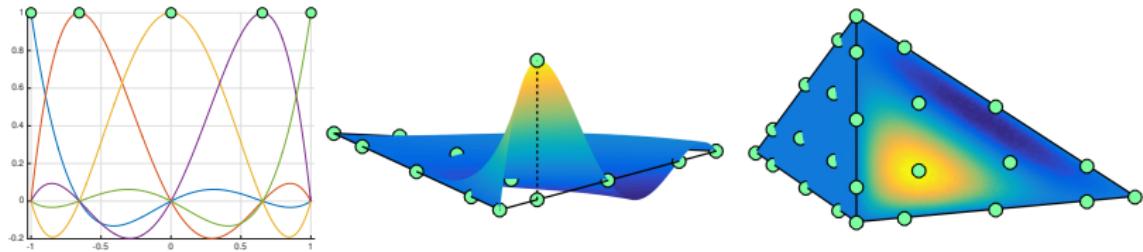


- Large **dense** matrices: $O(N^6)$ work per tet.
- High orders usually use tensor-product elements: $O(N^4)$ vs $O(N^6)$ cost, but less geometric flexibility.
- Idea: choose basis such that matrices are **sparse**.

WADG runtimes for 50 timesteps, 98304 elements.

BBDG: Bernstein-Bezier DG methods

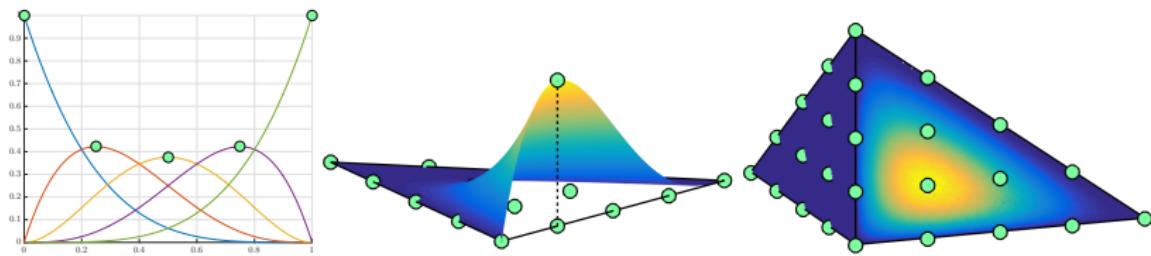
- Nodal DG: $O(N^6)$ cost in 3D vs $O(N^3)$ degrees of freedom.
- Switch to Bernstein basis: sparse and structured matrices.
- Optimal $O(N^3)$ application of differentiation and lifting matrices.



Nodal bases in one, two, and three dimensions.

BBDG: Bernstein-Bezier DG methods

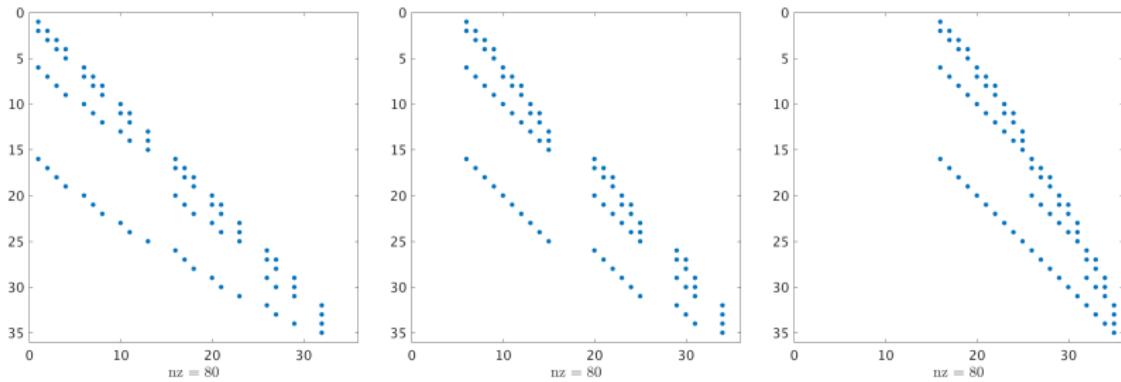
- Nodal DG: $O(N^6)$ cost in 3D vs $O(N^3)$ degrees of freedom.
- Switch to Bernstein basis: sparse and structured matrices.
- Optimal $O(N^3)$ application of differentiation and lifting matrices.



Bernstein bases in one, two, and three dimensions.

BBDG: Bernstein-Bezier DG methods

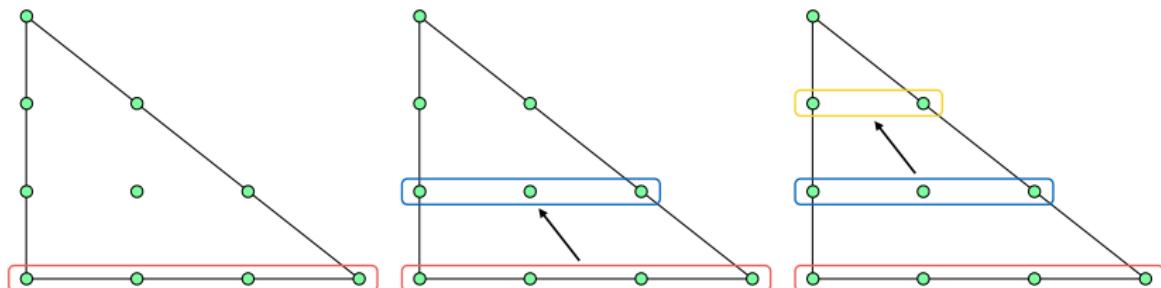
- Nodal DG: $O(N^6)$ cost in 3D vs $O(N^3)$ degrees of freedom.
- Switch to Bernstein basis: sparse and structured matrices.
- Optimal $O(N^3)$ application of differentiation and lifting matrices.



Sparse Bernstein differentiation matrices for the reference tetrahedron.

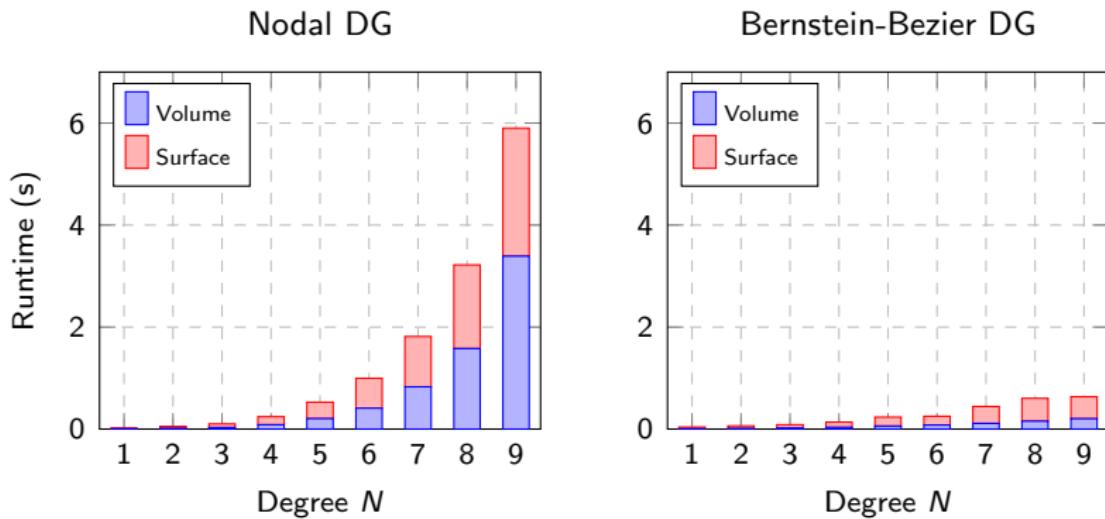
BBDG: Bernstein-Bezier DG methods

- Nodal DG: $O(N^6)$ cost in 3D vs $O(N^3)$ degrees of freedom.
- Switch to Bernstein basis: sparse and structured matrices.
- Optimal $O(N^3)$ application of differentiation and lifting matrices.



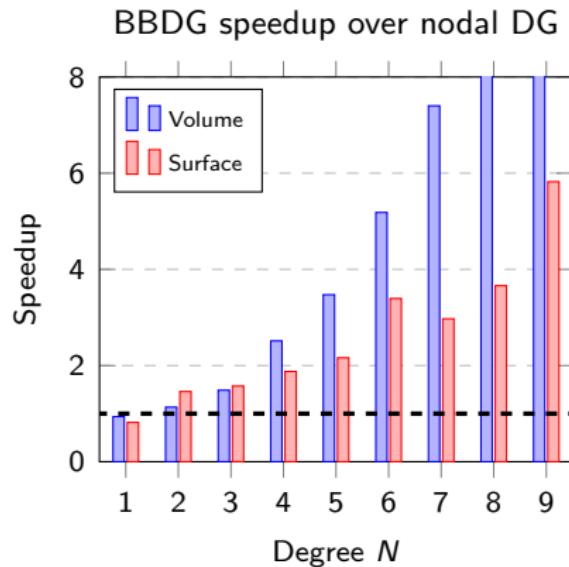
Optimal $O(N^3)$ complexity “slice-by-slice” application of Bernstein lift.

BBDG: efficient volume, surface kernels



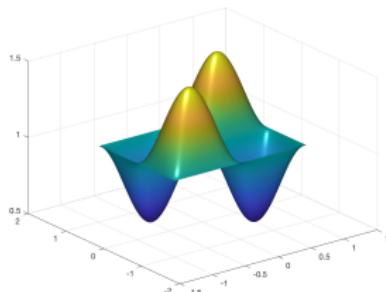
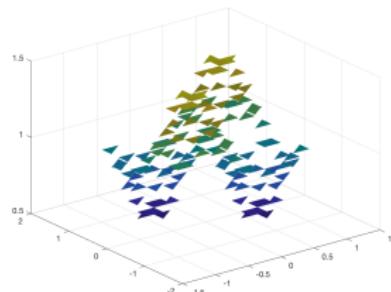
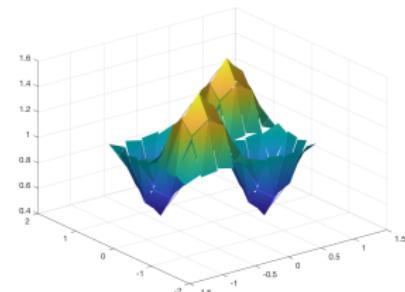
$$\underbrace{\frac{d\mathbf{u}}{dt}}_{\text{Update kernel}} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f}_{\text{Surface kernel}} (\text{flux}), \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

BBDG: efficient volume, surface kernels



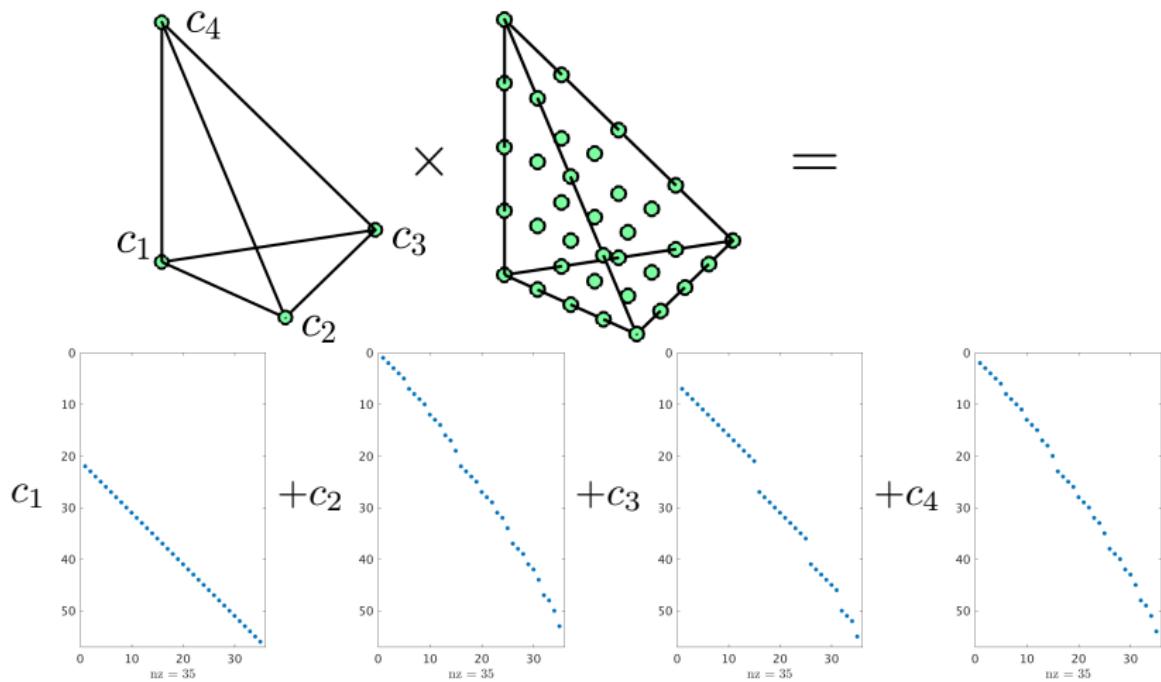
$$\underbrace{\frac{d\mathbf{u}}{dt}}_{\text{Update kernel}} = \underbrace{\mathbf{D}_x \mathbf{u}}_{\text{Volume kernel}} + \underbrace{\sum_{\text{faces}} \mathbf{L}_f}_{\text{Surface kernel}} (\text{flux}), \quad \mathbf{L}_f = \mathbf{M}^{-1} \mathbf{M}_f.$$

BBWADG: polynomial multiplication and projection

(a) Exact c^2 (b) $M = 0$ approximation(c) $M = 1$ approximation

- WADG: can reuse fast Bernstein volume and surface kernels.
- $O(N^6)$ update kernel: \mathbf{V}_q interpolates $u(\mathbf{x})$ to quadrature points, scale by $c^2(\mathbf{x})$ at quadrature points, apply \mathbf{P}_q to project back to P^N .
- New approach: approx. $c^2(\mathbf{x})$ with degree M polynomial, use fast Bernstein algorithms for polynomial multiplication and projection.

Fast Bernstein polynomial multiplication



Bernstein polynomial multiplication: for fixed M , $O(N^3)$ complexity.

Fast Bernstein polynomial projection

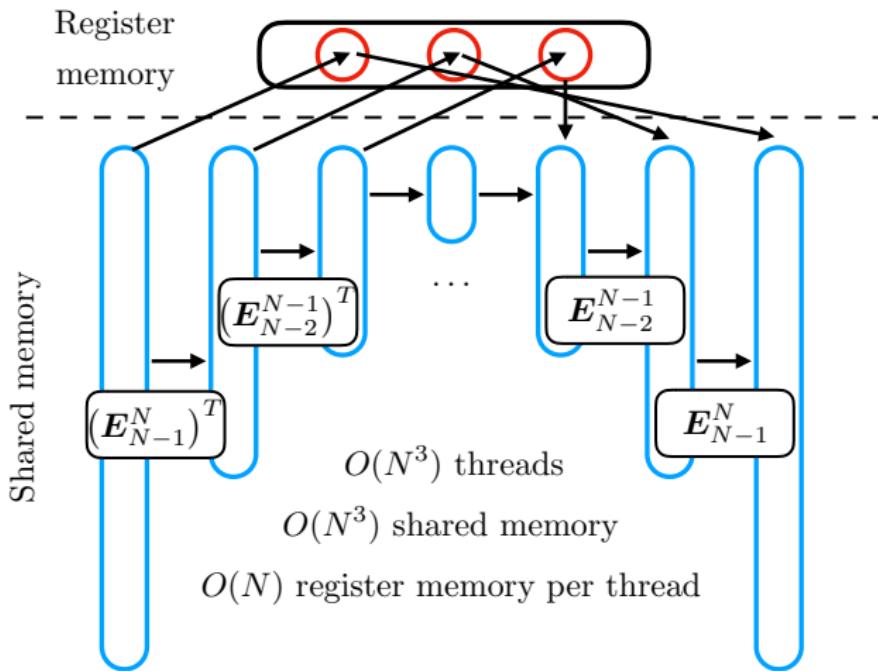
- Given $c^2(\mathbf{x})u(\mathbf{x})$ as a degree $(N + M)$ polynomial, apply L^2 projection matrix \mathbf{P}_N^{N+M} to reduce to degree N .
- Polynomial L^2 projection matrix \mathbf{P}_N^{N+M} under Bernstein basis:

$$\mathbf{P}_N^{N+M} = \underbrace{\sum_{j=0}^N c_j \mathbf{E}_{N-j}^N \left(\mathbf{E}_{N-j}^N \right)^T \left(\mathbf{E}_N^{N+M} \right)^T}_{\tilde{\mathbf{P}}_N}$$

- “Telescoping” form of $\tilde{\mathbf{P}}_N$: $O(N^4)$ complexity, more GPU-friendly.

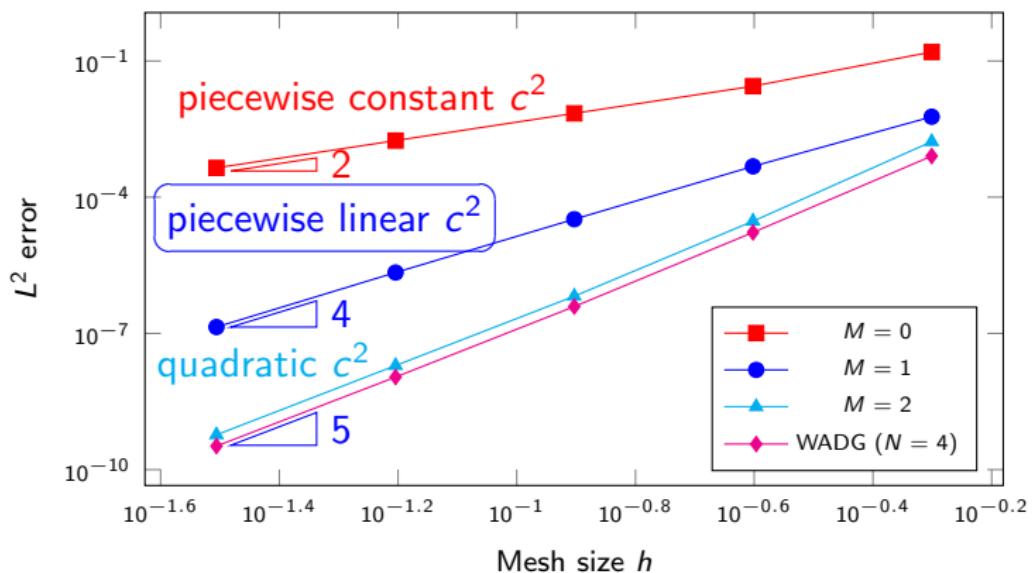
$$\left(c_0 \mathbf{I} + \mathbf{E}_{N-1}^N \left(c_1 \mathbf{I} + \mathbf{E}_{N-2}^{N-1} \left(c_2 \mathbf{I} + \cdots \right) \left(\mathbf{E}_{N-2}^{N-1} \right)^T \right) \left(\mathbf{E}_{N-1}^N \right)^T \right)$$

Sketch of GPU algorithm for \tilde{P}_N



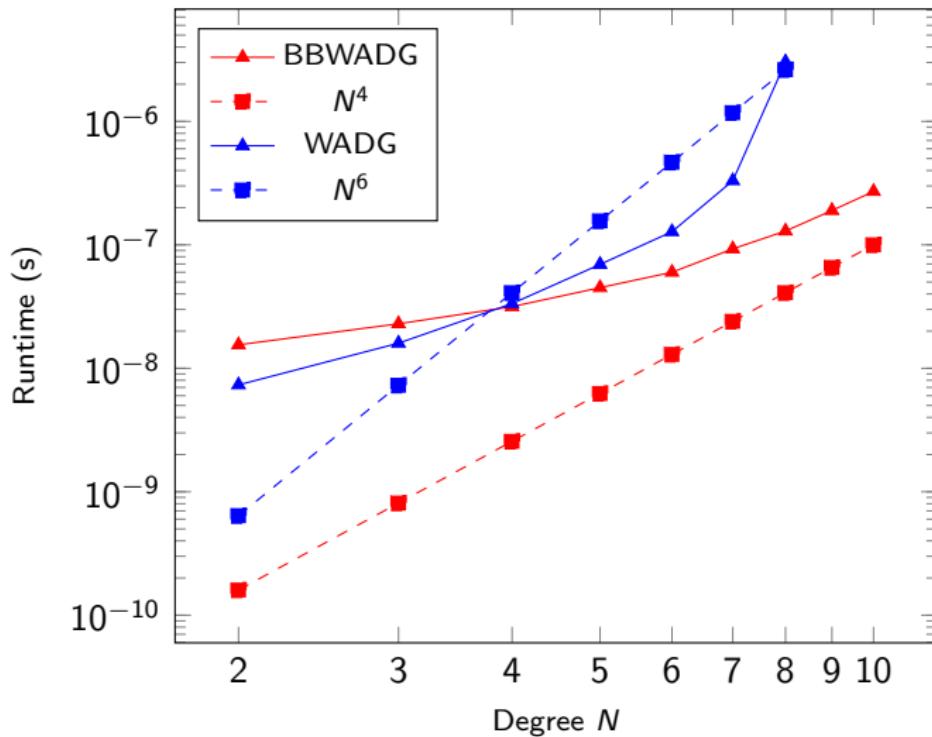
$$\left(c_0 \mathbf{I} + \mathbf{E}_{N-1}^N \left(c_1 \mathbf{I} + \mathbf{E}_{N-2}^{N-1} \left(c_2 \mathbf{I} + \cdots \right) \left(\mathbf{E}_{N-2}^{N-1} \right)^T \right) \left(\mathbf{E}_{N-1}^N \right)^T \right)$$

BBWADG: approximating c^2 and accuracy



Approximating smooth $c^2(x)$ using L^2 projection:
 $O(h^2)$ for $M = 0$, $O(h^4)$ for $M = 1$, $O(h^{M+3})$ for $0 < M \leq N - 2$.

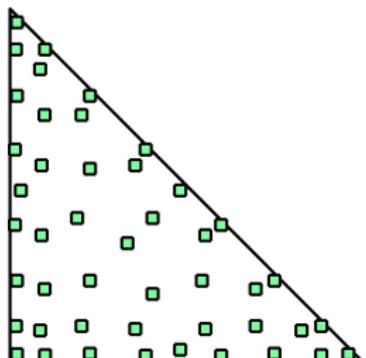
BBWADG: computational runtime (acoustics)

Update kernel for $M = 1$: runtime per element

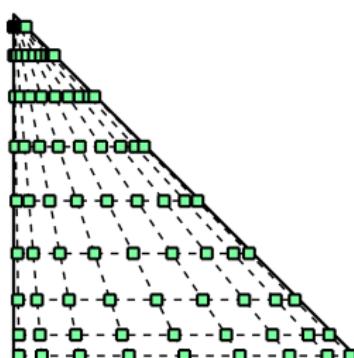
BBWADG: update kernel speedup over WADG (acoustics)

	$N = 3$	$N = 4$	$N = 5$	$N = 6$	$N = 7$	$N = 8$
WADG	1.60e-8	3.34e-8	6.94e-8	1.28e-7	3.31e-7	3.03e-6
BBWADG	2.20e-8	3.30e-8	4.42e-8	6.01e-8	9.46e-8	1.31e-7
Speedup	0.7260	1.0127	1.5706	2.1258	3.4938	23.1591

For $N \geq 8$, quadrature (and WADG) becomes much more expensive.



(a) $N = 7$ quadrature



(b) $N = 8$ quadrature

Summary and acknowledgements

- Weight-adjusted DG: stability and efficiency for heterogeneous media.
- BBWADG: improved complexity for approximate wavespeeds.
- This work is supported by the National Science Foundation under DMS-1712639 and DMS-1719818.

Thank you! Questions?



-
- Chan, Hewett, Warburton. 2016. Weight-adjusted DG methods: wave propagation in heterogeneous media (SISC).
Chan 2017. Weight-adjusted DG methods: matrix-valued weights and elastic wave prop. in heterogeneous media (IJNME).
Chan, Warburton 2015. GPU-accelerated Bernstein-Bezier discontinuous Galerkin methods for wave propagation (SISC).